

XYZ Platform – AWS Architecture Overview

A scalable, secure, and highly available cloud architecture designed

1. Architecture Summary

- **Deployment Model:** Multi-AZ (Availability Zone) for fault tolerance.
- **Tiered Design:**
 - **Web Tier:** User-facing layer (frontend).
 - **App Tier:** Business logic & microservices.
 - **Data Tier:** Databases & storage.
- **Scalability:** Auto-scaling, serverless components, and managed services.

2. Core Components

2.1 Networking & Subnets

- **VPC (Virtual Private Cloud):** Isolated network environment.
 - **Public Subnets (2+):**
 - NAT Gateways (for private instance internet access).
 - Load Balancers.
 - **Private Subnets (3+ per AZ):**
 - **Web Tier:** Frontend servers (e.g., EC2, ECS).
 - **App Tier:** Backend services (e.g., EC2, Lambda, EKS).
 - **Data Tier:** Databases (e.g., RDS, DynamoDB).

2.2 Traffic Routing & Load Balancing

- **DNS & Routing:**
 - **Amazon Route 53:** Domain management & latency-based routing.
- **Security & Traffic Filtering:**
 - **AWS WAF:** Protects against SQLi, XSS, DDoS.
 - **AWS Shield:** DDoS protection for ALB/CloudFront.
- **Content Delivery:**

- **CloudFront:** CDN for static/media content.
- **ALB (Application Load Balancer):** Distributes traffic to web/app tiers.

2.3 Compute Layer

- **Web Tier:**
 - EC2 Auto Scaling Groups (ASG) or Containers (ECS/EKS).
- **App Tier:**
 - Microservices (Lambda, ECS Fargate).
 - Event-driven processing (SQS, EventBridge).
- **Serverless Options:**
 - API Gateway + Lambda for APIs.

2.4 Data Layer

- **Primary Database:**
 - **Amazon RDS (Multi-AZ):** For relational data (PostgreSQL/MySQL).
 - **DynamoDB:** NoSQL for high-speed queries.
- **Caching:**
 - **ElastiCache (Redis/Memcached):** Session/store caching.
- **File Storage:**
 - **Amazon S3:** Static assets, user uploads.
 - **EFS:** Shared file storage for EC2.

2.5 Security

- **Identity & Access:**
 - **IAM:** Least-privilege roles for EC2/Lambda.
 - **Cognito:** User authentication.
- **Network Security:**
 - **Security Groups (SGs):** Tier-specific traffic rules.
 - **NACLs (Network ACLs):** Subnet-level firewall.

- **Encryption:**
 - KMS for data-at-rest, TLS for in-transit.

2.6 Monitoring & Operations

- **Observability:**
 - **CloudWatch:** Logs, metrics, alarms.
 - **X-Ray:** Distributed tracing.
- **Alerts:**
 - **SNS:** Notifications for failures/thresholds.
- **CI/CD:**
 - **CodePipeline/CodeDeploy:** Automated deployments.

3. High Availability (HA) & Disaster Recovery (DR)

- **Multi-AZ Deployments:** RDS, EC2 ASG.
- **Backups:**
 - RDS snapshots, S3 versioning.
- **Cross-Region Replication:**
 - S3 CRR, DynamoDB Global Tables.

4. Cost Optimization

- **Reserved Instances:** For steady-state workloads.
- **Spot Instances:** For fault-tolerant tasks.
- **Auto Scaling:** Scale down during off-peak.

5. Diagram

